

# Parallel and Distributed Computing

## Assignment 1 – OpenMP Programming

Ramyak Sharma

## Introduction

This assignment evaluates performance gains obtained using OpenMP parallel programming. Three computational problems were parallelized and execution times were compared for varying thread counts.

## 1 Question 1: DAXPY Loop

### Problem Statement

The DAXPY operation performs

$$X[i] = a \times X[i] + Y[i]$$

on vectors of size  $2^{16}$ . Execution time and speedup were measured while increasing thread count.

### Implementation

Sequential execution was first measured. The loop was then parallelized using OpenMP and executed with thread counts ranging from 2 to 12.

### Results

Execution statistics from the program are shown below.

```

Sequential Time = 5.8e-05 seconds

Threads = 2  Time = 0.000121  Speedup = 0.479338
Threads = 3  Time = 9.40003e-05  Speedup = 0.617019
Threads = 4  Time = 0.000105  Speedup = 0.55238
Threads = 5  Time = 9.00002e-05  Speedup = 0.644443
Threads = 6  Time = 0.000112  Speedup = 0.517857
Threads = 7  Time = 9.99998e-05  Speedup = 0.580001
Threads = 8  Time = 9.3e-05  Speedup = 0.623656
Threads = 9  Time = 0.000115  Speedup = 0.504349
Threads = 10  Time = 0.000112  Speedup = 0.517857
Threads = 11  Time = 9.59998e-05  Speedup = 0.604168
Threads = 12  Time = 0.000107001  Speedup = 0.542053

```

Figure 1: DAXPY Execution Statistics

## Observation

Sequential execution time is extremely small. Parallel execution introduces thread creation and synchronization overhead, causing parallel runtimes to become larger than sequential runtime.

Consequently, speedup remains below 1 for all thread counts. The DAXPY computation is memory-bound, meaning threads compete for memory bandwidth rather than computation time.

## Conclusion

Parallelization does not improve performance for this workload due to low computation per iteration and memory access limitations.

## 2 Question 2: Matrix Multiplication

### Problem Statement

Matrix multiplication of large matrices was parallelized using two methods:

- 1D threading
- 2D threading

## Implementation

Sequential matrix multiplication was implemented first. Two parallel strategies were tested:

- 1D threading parallelizes rows so each thread computes separate output rows.
- 2D threading distributes work across rows and columns simultaneously.

Execution time was measured across multiple thread counts.

## Results

### 1D Threading

```
Sequential Time = 0.097889 seconds

1D Parallel
Threads = 2  Time = 0.044314  Speedup = 2.20899
Threads = 3  Time = 0.030118  Speedup = 3.25018
Threads = 4  Time = 0.02787   Speedup = 3.51234
Threads = 5  Time = 0.025857  Speedup = 3.78578
Threads = 6  Time = 0.025696  Speedup = 3.8095
Threads = 7  Time = 0.021414  Speedup = 4.57126
Threads = 8  Time = 0.023641  Speedup = 4.14065
Threads = 9  Time = 0.023697  Speedup = 4.13086
Threads = 10  Time = 0.023661  Speedup = 4.13715
Threads = 11  Time = 0.021332  Speedup = 4.58883
Threads = 12  Time = 0.020701  Speedup = 4.72871
```

Figure 2: Matrix Multiplication using 1D Threading

### 2D Threading

```

2D Parallel
Threads = 2 Time = 0.135075 Speedup = 0.724701
Threads = 3 Time = 0.09266 Speedup = 1.05643
Threads = 4 Time = 0.070924 Speedup = 1.3802
Threads = 5 Time = 0.070377 Speedup = 1.39092
Threads = 6 Time = 0.060459 Speedup = 1.6191
Threads = 7 Time = 0.053086 Speedup = 1.84397
Threads = 8 Time = 0.053182 Speedup = 1.84064
Threads = 9 Time = 0.054356 Speedup = 1.80089
Threads = 10 Time = 0.05256 Speedup = 1.86242
Threads = 11 Time = 0.05691 Speedup = 1.72007
Threads = 12 Time = 0.055134 Speedup = 1.77547

```

Figure 3: Matrix Multiplication using 2D Threading

## Observation

1D threading provides strong speedup, reaching approximately  $4.7\times$  improvement at 12 threads. Work distribution across rows minimizes scheduling overhead.

2D threading improves load distribution but introduces extra scheduling overhead, resulting in smaller speedups compared to 1D threading.

Performance scaling slows after several threads due to memory bandwidth and cache limitations.

## Conclusion

Matrix multiplication benefits significantly from parallelization. 1D threading proved more efficient on the tested hardware due to lower overhead.

## 3 Question 3: Calculation of $\pi$

### Problem Statement

The value of  $\pi$  is computed using numerical integration:

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$

## Implementation

The computation loop was parallelized using OpenMP reduction to combine partial sums computed by threads.

## Results

```
Sequential:  
Pi = 3.14159  
Time = 0.110957  
  
Threads=2 Pi=3.14159 Time=0.049051 Speedup=2.26207  
Threads=3 Pi=3.14159 Time=0.033675 Speedup=3.29494  
Threads=4 Pi=3.14159 Time=0.02755 Speedup=4.02748  
Threads=5 Pi=3.14159 Time=0.026384 Speedup=4.20547  
Threads=6 Pi=3.14159 Time=0.024939 Speedup=4.44914  
Threads=7 Pi=3.14159 Time=0.020178 Speedup=5.49891  
Threads=8 Pi=3.14159 Time=0.022931 Speedup=4.83873  
Threads=9 Pi=3.14159 Time=0.023826 Speedup=4.65697  
Threads=10 Pi=3.14159 Time=0.022785 Speedup=4.86974  
Threads=11 Pi=3.14159 Time=0.021682 Speedup=5.11747  
Threads=12 Pi=3.14159 Time=0.021389 Speedup=5.18757
```

Figure 4: Parallel  $\pi$  Calculation Statistics

## Observation

Parallel execution shows strong improvement, reaching speedups above  $5\times$  around 7–12 threads. Performance variation occurs due to scheduling and memory contention effects.

Reduction overhead is minimal compared to computation cost, allowing efficient scaling.

## Conclusion

Parallel reduction effectively accelerates numerical integration. Optimal speedup occurs when thread count matches hardware capabilities.

## Overall Conclusion

OpenMP significantly improves performance for compute-intensive tasks such as matrix multiplication and numerical integration. However, memory-bound operations like

DAXPY may not benefit due to bandwidth limitations and thread overhead. Optimal thread count typically corresponds to available hardware cores.